

CSE 5311 Notes 8: Minimum Spanning Trees

(Last updated 2/18/17 4:10 PM)

CLRS, Chapter 23

CONCEPTS

Given a weighted, connected, undirected graph, find a minimum (total) weight free tree connecting the vertices. (AKA bottleneck shortest path tree)

Cut Property: Suppose S and T partition V such that

1. $S \cap T = \emptyset$
2. $S \cup T = V$
3. $|S| > 0$ and $|T| > 0$

then there is some MST that includes a minimum weight edge $\{s, t\}$ with $s \in S$ and $t \in T$.


Proof:

Suppose there is a partition with a minimum weight edge $\{s, t\}$.

A spanning tree without $\{s, t\}$ must still have a path between s and t .

Since $s \in S$ and $t \in T$, there must be at least one edge $\{x, y\}$ on this path with $x \in S$ and $y \in T$.

By removing $\{x, y\}$ and including $\{s, t\}$, a spanning tree whose total weight is no larger is obtained. •••



Cycle Property: Suppose a given spanning tree does not include the edge $\{u, v\}$. If the weight of $\{u, v\}$ is no larger than the weight of an edge $\{x, y\}$ on the unique spanning tree path between u and v , then replacing $\{x, y\}$ with $\{u, v\}$ yields a spanning tree whose weight does not exceed that of the original spanning tree.

Proof: Including $\{u, v\}$ into the spanning tree introduces a cycle, but removing $\{x, y\}$ will remove the cycle to yield a modified tree whose weight is no larger.

Does not directly suggest an algorithm, but all algorithms avoid including an edge that violates.

Prove or give counterexample:

The MST path between two vertices is a shortest path.

True or False?

Choosing the $|V| - 1$ edges with smallest weights gives a MST.

Fill in the blank:

Multiple MSTs occur only if _____.


MODIFIED REACHABILITY CONDITION

Towards an algorithm:

1. Assume unique edge weights (easily forced by lexicographically breaking ties).
2. Consider all (cycle-free) paths between some pair of vertices.
3. Consider the maximum weight edge on each path.
4. The edge that is the minimum of the “maximums” *must* be included in the MST.

Any edge that is never a “must” is not in the MST.

Example:



Original Adjacency Matrix							Unchanged Entries ($\boxed{}$) are MST						
	A	B	C	D	E	F		A	B	C	D	E	F
A	∞	1	2	4	∞	∞	A	1	$\boxed{1}$	$\boxed{2}$	$\boxed{4}$	5	7
B	1	∞	3	∞	∞	∞	B	$\boxed{1}$	1	2	4	5	7
C	2	3	∞	∞	5	∞	C	$\boxed{2}$	2	2	4	$\boxed{5}$	7
D	4	∞	∞	∞	6	7	D	$\boxed{4}$	4	4	4	5	$\boxed{7}$
E	∞	∞	5	6	∞	8	E	5	5	$\boxed{5}$	5	5	7
F	∞	∞	∞	7	8	∞	F	7	7	7	$\boxed{7}$	7	7



Implementation - Based on Warshall's Algorithm

1. Directed reachability - existence of path:

```

for (j=0; j<V; j++)
  for (i=0; i<V; i++)
    if (A[i][j])
      for (k=0; k<V; k++)
        if (A[j][k])
          A[i][k]=1;


```




Correctness proof is by math. induction. (CSE 2320 Notes 16.C)

Successor Matrix (CLRS uses predecessor)


7-11 directions:



Initialize:



Warshall Matrix Update:



```

for (j=0; j<V; j++)
    for (i=0; i<V; i++)
        if (s[i][j] != (-1))
            for (k=0; k<V; k++)
                if (s[i][k] == (-1) && s[j][k] != (-1))
                    s[i][k] = s[i][j];

```

2. All-pairs shortest paths - Floyd-Warshall

```

for (j=0; j<V; j++)
    for (i=0; i<V; i++)
        if (dist[i][j] < 999)
            for (k=0; k<V; k++)
                {
                    newDist = dist[i][j] + dist[j][k];
                    if (newDist < dist[i][k])
                        {
                            dist[i][k] = newDist;
                            s[i][k] = s[i][j];
                        }
                }

```

3. Minimum spanning tree: <http://ranger.uta.edu/~weems/NOTES5311/MSTWarshall.c>

```

// MST based on Warshall's algorithm (Maggs & Plotkin,
//   Information Processing Letters 26, 25 Jan 1988, 291-293)
// 3/6/03 BPW

// Modified 7/15/04 to make more robust, especially edges with same weight

#include <stdio.h>

#define maxSize (20)

struct edge {
    int weight, smallLabel, largeLabel;
};
typedef struct edge edgeType;

edgeType min(edgeType x, edgeType y)
{
    // Returns smaller-weighted edge, using lexicographic tie-breaker
    if (x.weight < y.weight)
        return x;
    if (x.weight > y.weight)
        return y;

    if (x.smallLabel < y.smallLabel)
        return x;
    if (x.smallLabel > y.smallLabel)
        return y;

    if (x.largeLabel < y.largeLabel)
        return x;
    return y;
}

```

```

edgeType max(edgeType x,edgeType y)
{
// Returns larger-weighted edge, using lexicographic tie-breaker
if (x.weight>y.weight)
    return x;
if (x.weight<y.weight)
    return y;

if (x.smallLabel>y.smallLabel)
    return x;
if (x.smallLabel<y.smallLabel)
    return y;

if (x.largeLabel>y.largeLabel)
    return x;
return y;
}

main()
{
int numVertices,numEdges, i, j, k;
edgeType matrix[maxSize][maxSize];
int count;

printf("enter # of vertices and edges: ");
fflush(stdout);
scanf("%d %d",&numVertices,&numEdges);
printf("enter undirected edges u v weight\n");
for (i=0;i<numVertices;i++)
    for (j=0;j<numVertices;j++)
    {
        matrix[i][j].weight=999;
        if (i<=j)
        {
            matrix[i][j].smallLabel=i;
            matrix[i][j].largeLabel=j;
        }
        else
        {
            matrix[i][j].smallLabel=j;
            matrix[i][j].largeLabel=i;
        }
    }
for (k=0;k<numEdges;k++)
{
    scanf("%d %d",&i,&j);
    scanf("%d",&matrix[i][j].weight);
    matrix[j][i].weight=matrix[i][j].weight;
}

printf("input matrix\n");
for (i=0;i<numVertices;i++)
{
    for (j=0;j<numVertices;j++)
        printf("%3d ",matrix[i][j].weight);
    printf("\n");
}
}

```

```
// MST by Warshall
for (k=0;k<numVertices;k++)
  for (i=0;i<numVertices;i++)
    for (j=0;j<numVertices;j++)
      matrix[i][j]=min(matrix[i][j],max(matrix[i][k],matrix[k][j]));

printf("output matrix\n");
for (i=0;i<numVertices;i++)
{
  for (j=0;j<numVertices;j++)
    printf("%3d(%3d,%3d) ",matrix[i][j].weight,matrix[i][j].smallLabel,
      matrix[i][j].largeLabel);
  printf("\n");
}

count=0;
for (i=0;i<numVertices;i++)
  for (j=i+1;j<numVertices;j++)
    if (matrix[i][j].weight<999 && i==matrix[i][j].smallLabel &&
      j==matrix[i][j].largeLabel)
    {
      count++;
      printf("%d %d %d\n",i,j,matrix[i][j].weight);
    }

if (count<numVertices-1)
  printf("Result is a spanning forest\n");
else if (count>=numVertices)
  printf("Error? . . . \n");
}
```


PRIM'S ALGORITHM

Outline:

1. Vertex set S: tree that grows to MST.
2. $T = V - S$: vertices not yet in tree.
3. Initialize S with arbitrary vertex.
4. Each step moves one vertex from T to S: the one with the minimum weight edge to an S vertex.

Different data structures lead to various performance characteristics.

Which edge does Prim's algorithm select next?



1. Maintains T-table that provides the closest vertex in S for each vertex in T.

Scans the list of the last vertex moved from T to S.

Place any vertex $x \in V$ in S.

$T = V - \{x\}$

for each $t \in T$

Initialize T-table entry with weight of $\{t, x\}$ (or ∞ if non-existent) and x as best-S-neighbor.

while $T \neq \emptyset$

Scan T-table entries for the minimum weight edge $\{t, \text{best-S-neighbor}[t]\}$

over all $t \in T$ and all $s \in S$.

Include edge $\{t, \text{best-S-neighbor}[t]\}$ in MST.

$T = T - \{t\}$

$S = S \cup \{t\}$


for each vertex x in adjacency list of t

if $x \in T$ and weight of $\{x, t\}$ is smaller than T-weight[x]

T-weight[x] = weight of $\{x, t\}$

best-S-neighbor[x] = t

What are the T-table contents before and after the next MST vertex is selected?



Analysis:

Initializing the T-table takes $\Theta(V)$.
 Scans of T-table entries contribute $\Theta(V^2)$.
 Traversals of adjacency lists contribute $\Theta(E)$.
 $\Theta(V^2 + E)$ overall worst-case.

2. Replace T-table by a heap.

The time for updating for best-S-neighbor increases, but the time for selection of the next vertex to move from T to S improves.

Place any vertex $x \in V$ in S.

$T = V - \{x\}$

for each $t \in T$

Load T-heap entry with weight (as the priority) of $\{t, x\}$ (or ∞ if non-existent) and x as best-S-neighbor

BUILD-MIN-HEAP(T-heap)

while $T \neq \emptyset$

Use HEAP-EXTRACT-MIN to obtain T-heap entry with the minimum weight edge over all $t \in T$ and all $s \in S$.

Include edge $\{t, \text{best-S-neighbor}[t]\}$ in MST.

$T = T - \{t\}$

$S = S \cup \{t\}$

for each vertex x in adjacency list of t

if $x \in T$ and weight of $\{x, t\}$ is smaller than T-weight[x]

T-weight[x] = weight of $\{x, t\}$

best-S-neighbor[x] = t

MIN-HEAP-DECREASE-KEY(T-heap)

Analysis for binary heap:

Initializing the T-heap takes $\Theta(V)$.
 Total cost for HEAP-EXTRACT-MINS is $\Theta(V \log V)$.
 Traversals of adjacency lists and MIN-HEAP-DECREASE-KEYS contribute $\Theta(E \log V)$.
 $\Theta(E \log V)$ overall worst-case, since $E > V$.

Analysis (amortized) for Fibonacci heap:

Initializing the T-heap takes $\Theta(V)$.
 Total cost for HEAP-EXTRACT-MINS is $\Theta(V \log V)$.
 Traversals of adjacency lists and MIN-HEAP-DECREASE-KEYS contribute $\Theta(E)$.
 $\Theta(E + V \log V)$ overall worst-case, since $E > V$.

Which version is the fastest?

Sparse ($E = O(V)$) Dense ($E = \Omega(V^2)$)

table	$\Theta(V^2 + E)$	$\Theta(V^2)$	$\Theta(V^2)$
binary heap	$\Theta(E \log V)$	$\Theta(V \log V)$	$\Theta(V^2 \log V)$
Fibonacci heap	$\Theta(E + V \log V)$	$\Theta(V \log V)$	$\Theta(V^2)$

Analysis also applies to Dijkstra's shortest path.

KRUSKAL'S ALGORITHM

(Discussed in Notes 7 as an application of UNION-FIND trees.)

<http://ranger.uta.edu/~weems/NOTES5311/kruskal.c>

. . .

main()

{

. . .

qsort(edgeTab, numEdges, sizeof(edgeType), weightAscending);

for (i=0; i<numEdges; i++)

{

 root1=find(edgeTab[i].tail);

 root2=find(edgeTab[i].head);

 if (root1==root2)

 printf("%d %d %d discarded\n", edgeTab[i].tail, edgeTab[i].head,
 edgeTab[i].weight);

 else

 {

 printf("%d %d %d included\n", edgeTab[i].tail, edgeTab[i].head,
 edgeTab[i].weight);

 MSTweight+=edgeTab[i].weight;

 makeEquivalent(root1, root2);

 }

}

if (numTrees!=1)

 printf("MST does not exist\n");

printf("Sum of weights of spanning edges %d\n", MSTweight);

}


Incremental sorting (e.g. QUICKSORT or HEAPSORT) may be used.

Can adapt to determine if MST is unique:

<http://ranger.uta.edu/~weems/NOTES5311/kruskalDup.c>

. . .

```
main()
{
    . . .
    numTrees=numVertices;
    qsort(edgeTab,numEdges,sizeof(edgeType),weightAscending);
    i=0;
    while (i<numEdges)
    {
        for (k=i;
            k<numEdges && edgeTab[k].weight==edgeTab[i].weight;
            k++)
            ;
        for (j=i;j<k;j++)
        {
            root1=find(edgeTab[j].tail);
            root2=find(edgeTab[j].head);
            if (root1==root2)
            {
                printf("%d %d %d discarded\n",edgeTab[j].tail,edgeTab[j].head,
                    edgeTab[j].weight);
                edgeTab[j].weight=(-1);
            }
        }
        for (j=i;j<k;j++)
            if (edgeTab[j].weight!=(-1))
            {
                root1=find(edgeTab[j].tail);
                root2=find(edgeTab[j].head);
                if (root1==root2)
                    printf("%d %d %d alternate\n",edgeTab[j].tail,edgeTab[j].head,
                        edgeTab[j].weight);
                else
                {
                    printf("%d %d %d included\n",edgeTab[j].tail,edgeTab[j].head,
                        edgeTab[j].weight);
                    makeEquivalent(root1,root2);
                }
            }
        i=k;
    }
    if (numTrees!=1)
        printf("MST does not exist\n");
}
```



BORUVKA'S ALGORITHM

Similar to Kruskal:

1. Initially, each vertex is a component.
2. Each component has a "best edge" to some other component.
3. Boruvka step:

For each best edge from a component x to a component y:




```

a = FIND(x)
b = FIND(y)
if a ≠ b
    UNION(a,b)
  
```

Worst-case: number of components decreases by at least half in each phase.

Gives $O(E \log V)$ time.

In some cases a cluster of several components may collapse:



<http://ranger.uta.edu/~weems/NOTES5311/boruvka.c>

```

. . .
main()
{
. . .
numTrees=numVertices; // Each vertex is initially in its own subtree
usefulEdges=numEdges; // An edge is useful if the two vertices are separate
while (numTrees>1 && usefulEdges>0)
{
    for (i=0;i<numVertices;i++)
        bestEdgeNum[i]=(-1);
    usefulEdges=0;
    for (i=0;i<numEdges;i++)
    {
        root1=find(edgeTab[i].tail);
        root2=find(edgeTab[i].head);
        if (root1==root2)
            printf("%d %d %d useless\n",edgeTab[i].tail,edgeTab[i].head,
                edgeTab[i].weight);
        else
        {
            usefulEdges++;
            if (bestEdgeNum[root1]==(-1)
                || edgeTab[bestEdgeNum[root1]].weight>edgeTab[i].weight)
                bestEdgeNum[root1]=i; // Have a new best edge from this component

            if (bestEdgeNum[root2]==(-1)
                || edgeTab[bestEdgeNum[root2]].weight>edgeTab[i].weight)
                bestEdgeNum[root2]=i; // Have a new best edge from this component
        }
    }
    for (i=0;i<numVertices;i++)
        if (bestEdgeNum[i]!=(-1))
        {
            root1=find(edgeTab[bestEdgeNum[i]].tail);
            root2=find(edgeTab[bestEdgeNum[i]].head);
            if (root1==root2)
                continue; // This round has already connected these components.
            MSTweight+=edgeTab[bestEdgeNum[i]].weight;
            printf("%d %d %d included in MST\n",
                edgeTab[bestEdgeNum[i]].tail,edgeTab[bestEdgeNum[i]].head,
                edgeTab[bestEdgeNum[i]].weight);
            makeEquivalent(root1,root2);
        }
    printf("numTrees is %d\n",numTrees);
}
if (numTrees!=1)
    printf("MST does not exist\n");
printf("Sum of weights of spanning edges %d\n",MSTweight);
}

```

ASIDE - COUNTING SPANNING TREES

https://en.wikipedia.org/wiki/Kirchhoff's_theorem